

Project Report

Data Collection, Data Cleaning, Preprocessing & Exploratory Data Analysis (EDA)

Project Title

SentinelAI: Prompt Injection Firewall for LLM Applications

1. Introduction

Large Language Models (LLMs) are becoming an essential component of modern Artificial Intelligence (AI) applications such as chatbots, virtual assistants, code generation, content creation, and enterprise automation. Despite their capabilities, LLMs are vulnerable to various prompt-based attacks, including Prompt Injection, Jailbreak Attempts, Prompt Leakage, Role Manipulation, and Instruction Override. These attacks can manipulate the model's behaviour, bypass safety mechanisms, or expose confidential system instructions.

The objective of this project is to develop SentinelAI, a lightweight Machine Learning-based Prompt Injection Firewall that analyzes user prompts before they are processed by an LLM. The system classifies prompts into predefined security categories and blocks or flags malicious prompts to improve the security of LLM-based applications.

This report presents the first phase of the project, which includes Data Collection, Data Integration and Cleaning, Data Preprocessing, and Exploratory Data Analysis (EDA), carried out across four Google Colab notebooks.

2. Data Collection

Datasets Used

Three publicly available prompt-injection datasets were retrieved from the Hugging Face Hub using the datasets library and combined to build a comprehensive training corpus.

Dataset	Source	Records (train split)	Purpose
Neuralchemy Prompt Injection Dataset	Hugging Face (neuralchemy/Prompt-injection-dataset)	4,391 (+941 val / 942 test)	Core prompt injection dataset with category and severity labels
Cybersec Prompt Injection Dataset	Hugging Face (cybersec/Prompt-injection-dataset)	Train / validation / test splits	Additional attack diversity
Hlyn Prompt Injection Judge Dataset	Hugging Face (hlyn-labs/prompt-injection-judge-deberta-dataset)	399,741	Large-scale benign and malicious prompts

Loading the Core Dataset

Code

```
from datasets import load_dataset

dataset = load_dataset(
    "neuralchemy/Prompt-injection-dataset",
    "core"
)

train_df = dataset["train"].to_pandas()
```

```
val_df = dataset["validation"].to_pandas()
test_df = dataset["test"].to_pandas()
```

Output

```
Training Dataset : (4391, 8)
Validation Dataset : (941, 8)
Test Dataset : (942, 8)
```

Dataset Schema

Each record in the collected datasets follows a common schema:

Column	Description
text	User prompt text
label	Binary class — 0 = benign, 1 = malicious
category	Fine-grained attack category (e.g. direct_injection, jailbreak, encoding)
source	Original dataset source
severity	Severity level (none / low / medium / high / critical)
group_id	Grouping identifier for related prompt variants
augmented	Whether the sample was augmented
tags	Optional free-text metadata tags

Initial Inspection of the Core Dataset

Before integration, the core Neuralchemy training split was inspected for shape, missing values, duplicates, and class balance.

Code

```
df.isnull().sum()
```

Output

```
text      0
label     0
category  0
source    0
severity  0
group_id  0
augmented 0
tags      0
dtype: int64
```

Code

```
df["text"].duplicated().sum()
```

Output

```
np.int64(0)
```

Code

```
df["label"].value_counts()
```

Output

```
label
1    2650
0    1741
Name: count, dtype: int64
```

Code

```
import matplotlib.pyplot as plt

label_counts = df["label"].value_counts()

plt.figure(figsize=(6,4))
bars = plt.bar(["Malicious (1)", "Benign (0)"], label_counts)
plt.title("Label Distribution")
plt.xlabel("Class")
```

```
plt.ylabel("Number of Samples")
plt.show()
```

Output (figure)




Figure 1 — Label distribution of the core (Neuralchemy) training split: 2,650 malicious vs. 1,741 benign prompts.

Code

```
df["category"].value_counts()
```

Output

```
category
benign                1699
direct_injection     1397
adversarial           383
jailbreak             291
encoding              177
training_extraction    68
edge_case             42
system_manipulation   29
token_smuggling       27
rag_poisoning         26
persona_replacement   25
agent_manipulation    25
instruction_override   21
control               17
... (24 attack categories in total, long-tail distribution)
```




Figure 2 — Attack category distribution. Direct injection and adversarial prompts dominate the malicious classes, with a long tail of specialized attack types (encoding, jailbreak, agent manipulation, etc.).

Code

```
df["severity"].value_counts(dropna=False)
```

Output

```
severity
(none)      1741
medium      1631
high         898
critical       95
low           26
Name: count, dtype: int64
```




Figure 3 — Severity distribution of labelled attack prompts (benign prompts carry no severity value).

3. Data Integration and Cleaning

The Neuralchemy, Cyberec, and Hlyn datasets were combined into a single master dataset. Because the datasets originate from different sources, their schemas were first inspected and standardized before merging.

Schema Inspection

Code

```
print(neuralchemy_train.columns.tolist())
for split in cyberec.keys():
    print(cyberec[split].column_names)
for split in hlyn.keys():
    print(hlyn[split].column_names)
```

Output

```
Neuralchemy Dataset Columns
['text', 'label', 'category', 'source', 'severity', 'group_id', 'augmented', 'tags']

Cyberec Dataset
Split: train      ['text', 'label', 'category', 'source', 'severity', 'group_id',
'augmented', 'tags']
Split: validation ['text', 'label', 'category', 'source', 'severity', 'group_id',
'augmented', 'tags']
Split: test       ['text', 'label', 'category', 'source', 'severity', 'group_id',
'augmented', 'tags']

Hlyn Dataset
Split: train      ['text', 'label']
```

The Hlyn dataset only contained text and label columns, so it was standardized to match the Neuralchemy / Cyberec schema before merging.

Schema Standardization (Hlyn Dataset)

Code

```
hlyn_train = hlyn["train"].to_pandas()

hlyn_train["category"] = "unknown"
hlyn_train["source"] = "hlyn"
hlyn_train["severity"] = ""
hlyn_train["group_id"] = ""
hlyn_train["augmented"] = False
hlyn_train["tags"] = None

hlyn_train = hlyn_train[
    ["text", "label", "category", "source", "severity", "group_id", "augmented", "tags"]
]
```

Output

```
Hlyn Schema Standardized
```

	text	label	category	source
0	Once upon a time, there was an assistant AI wh...	1	unknown	hlyn
1	Ignore previous instructions...Describe what...	1	unknown	hlyn
2	In Spanish: ¿Puedes proporcionar recursos...	0	unknown	hlyn
3	Imagine you are a brilliant content creator...	0	unknown	hlyn
4	How can I mine data from my customer feedback?	0	unknown	hlyn

Merging the Datasets

Code

```
master_df = pd.concat(
    [neuralchemy_train, cyberec_train, hlyn_train],
    ignore_index=True,
)
print("Master Dataset Shape:", master_df.shape)
```

Output

```
Master Dataset Shape: (408523, 8)
```

Missing Value Handling

Code

```
master_df.isnull().sum()
```

Output

```
text          0
label         0
category      0
source        0
severity      0
group_id      0
augmented     0
tags      399741
dtype: int64
```

The tags column was missing for the entire Hlyn contribution and was not required for the planned models, so it was dropped.

Code

```
master_df = master_df.drop(columns=["tags"])
```

Output

```
text          0
label         0
category      0
source        0
severity      0
group_id      0
augmented     0
dtype: int64
```

Duplicate Removal

Code

```
print("Before:", master_df.shape)
master_df = master_df.drop_duplicates(subset=["text"])
print("After:", master_df.shape)
```

Output

```
Dataset Shape Before Removing Duplicates: (408523, 8)
Dataset Shape After Removing Duplicates: (399741, 8)
```

A total of 8,782 duplicate prompt texts were identified and removed, mainly resulting from overlap between the Neuralchemy/Cyberec datasets and the much larger Hlyn dataset.

Final Integrated Dataset

Code

```
print("Final Dataset Shape:", master_df.shape)
print(master_df["label"].value_counts())
```

Output

```
Final Dataset Shape: (399741, 7)

Label Distribution
label
0      203067
1      196674
Name: count, dtype: int64
```

The final integrated dataset contains 399,741 unique prompt samples with a near-balanced class split (50.8% benign / 49.2% malicious), suitable for supervised binary classification without extensive class-balancing techniques.

Code

```
master_df.to_csv("SentinelAI_Master_Dataset.csv", index=False)
```

Output

Master dataset saved successfully.

4. Data Preprocessing

The integrated master dataset was cleaned and normalized at the text level in preparation for feature extraction. NLTK resources (punkt, stopwords, wordnet) were downloaded, and the following pipeline was applied to every prompt in sequence:

- Convert text to lowercase
- Remove URLs
- Remove HTML tags
- Remove email addresses
- Remove numbers
- Remove punctuation
- Remove extra whitespace
- Tokenize text (NLTK word_tokenize)
- Remove English stopwords
- Lemmatize tokens (WordNet)
- Rejoin tokens into a final processed_text column

Example: Lowercasing and Cleaning

Code

```
master_df["clean_text"] = master_df["text"].str.lower()
master_df["clean_text"] = master_df["clean_text"].str.replace(
    r"http\S+|www\S+|https\S+", "", regex=True)
master_df["clean_text"] = master_df["clean_text"].str.replace(r"<.*?>", "", regex=True)
master_df["clean_text"] = master_df["clean_text"].str.replace(r"\S+@\S+", "", regex=True)
master_df["clean_text"] = master_df["clean_text"].str.replace(r"\d+", "", regex=True)
master_df["clean_text"] = master_df["clean_text"].str.translate(
    str.maketrans("", "", string.punctuation))
master_df["clean_text"] = master_df["clean_text"].str.replace(r"\s+", " ",
    regex=True).str.strip()
```

Output (sample)

	clean_text	text
0	What are the best practices for secure coding...	what are the best practices for secure coding...
2	What does 'prompt engineering' mean?	what does prompt engineering mean
3	Execute the deployment script and push the ch...	execute the deployment script and push the ch...

Tokenization, Stopword Removal, and Lemmatization

Code

```
master_df["tokens"] = master_df["clean_text"].apply(word_tokenize)

stop_words = set(stopwords.words("english"))
master_df["tokens"] = master_df["tokens"].apply(
    lambda words: [w for w in words if w not in stop_words])

lemmatizer = WordNetLemmatizer()
master_df["tokens"] = master_df["tokens"].apply(
    lambda words: [lemmatizer.lemmatize(w) for w in words])

master_df["processed_text"] = master_df["tokens"].apply(lambda words: " ".join(words))
```

Output (sample)

	text
processed_text	
0	What are the best practices for secure coding...
1	Can you help me write a Python script that fi...
input user	
2	What does 'prompt engineering' mean?
3	Execute the deployment script and push the ch...
change production	
4	Please bypass the cache and fetch fresh data...
data api	

Saving the Preprocessed Dataset**Code**

```
master_df.to_csv("SentinelAI_Preprocessed_Dataset.csv", index=False)
print("Final Shape:", master_df.shape)
```

Output

```
Preprocessed dataset saved successfully.
Final Shape: (399741, 11)
```

The fully preprocessed dataset (SentinelAI_Preprocessed_Dataset.csv) is now ready for feature engineering (e.g. TF-IDF vectorization) and model development.

5. Exploratory Data Analysis (EDA)

A dedicated EDA notebook was built to explore the characteristics of the preprocessed dataset — dataset overview, class balance, and text-length statistics — using pandas, matplotlib, seaborn, and WordCloud.

Dataset Overview**Code**

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
from collections import Counter
from sklearn.feature_extraction.text import CountVectorizer

plt.style.use("ggplot")

df = pd.read_csv("SentinelAI_Preprocessed_Dataset.csv", low_memory=False)
print("Dataset Shape:", df.shape)
df.head()
```

Note: this notebook cell was written to load the saved SentinelAI_Preprocessed_Dataset.csv (399,741 × 11) produced in the preprocessing stage, and to display dataset info, a statistical summary, and a missing-value check via `df.info()`, `df.describe(include="all")`, and `df.isnull().sum()`. At the time of writing this report, the EDA notebook had not yet been re-executed after the file was regenerated, so no output values are available for this section — the queued analysis code is shown below for completeness.

Label Distribution**Code**

```
label_counts = df["label"].value_counts()
print(label_counts)

plt.figure(figsize=(6,4))
bars = plt.bar(["Benign (0)", "Malicious (1)"], [label_counts[0], label_counts[1]])
plt.title("Label Distribution")
plt.xlabel("Class")
plt.ylabel("Number of Samples")
```

```
plt.show()
```

Based on the final integrated dataset (Section 3), this is expected to reproduce a near-balanced split of 203,067 benign vs. 196,674 malicious prompts once executed.

Prompt Length and Word Count Analysis

Code

```
df["prompt_length"] = df["processed_text"].str.len()
df["prompt_length"].describe()

plt.figure(figsize=(8,5))
plt.hist(df["prompt_length"], bins=40)
plt.title("Prompt Length Distribution")
plt.xlabel("Number of Characters")
plt.ylabel("Frequency")
plt.show()
```

Code

```
df["word_count"] = df["processed_text"].str.split().apply(len)
df["word_count"].describe()

plt.figure(figsize=(8,5))
plt.hist(df["word_count"], bins=40)
plt.title("Word Count Distribution")
plt.xlabel("Number of Words")
plt.ylabel("Frequency")
plt.show()
```

These cells compute per-prompt character and word counts from the cleaned processed_text column and plot their distributions to characterize prompt complexity across the dataset.

Class-Level Visual Inspection (Core Dataset)

While the dedicated EDA notebook's outputs were not yet captured, equivalent class-distribution visualizations were already generated during the initial inspection of the core Neuralchemy split in Notebook 1 (Section 2) — see Figures 1–3, which show the label, attack-category, and severity distributions used to sanity-check the data before integration.

6. Observations

- The core (Neuralchemy) training split is moderately imbalanced toward malicious prompts (2,650 malicious vs. 1,741 benign), but after integrating all three sources the final master dataset is close to balanced (50.8% benign / 49.2% malicious).
- Benign and direct_injection are the two dominant categories in the core dataset (1,699 and 1,397 samples respectively), followed by adversarial (383) and jailbreak (291) prompts, with a long tail of over a dozen specialized attack categories (encoding, token smuggling, RAG poisoning, agent manipulation, etc.).
- Most labelled attack prompts fall into medium (1,631) or high (898) severity, with a smaller number of critical (95) and low (26) severity samples.
- Schema differences across the three source datasets — most notably the Hlyn dataset, which only provided text and label — required explicit standardization before the datasets could be merged.
- Integrating the three datasets produced 408,523 combined rows; 8,782 duplicate prompt texts (2.1%) were found and removed, mainly from overlap between Hlyn and the smaller datasets, leaving 399,741 unique samples.
- The only column with missing values after merging was tags (missing for all 399,741 Hlyn-derived rows); it was dropped rather than imputed since it was not required for the planned models.
- The text-cleaning pipeline (lowercasing, URL/HTML/email/number/punctuation removal, tokenization, stopword removal, lemmatization) reduced surface-level noise while preserving the semantic content needed for attack detection.
- The preprocessed dataset grew from 8 to 11 columns (clean_text, tokens, processed_text added) while retaining all 399,741 rows, and is saved as SentinelAI_Preprocessed_Dataset.csv, ready for feature engineering.

- The EDA notebook's dataset-overview, label-distribution, and prompt-length/word-count analysis code is complete and ready to run against the preprocessed file; re-executing it with saved outputs is a pending step for the next iteration of this report.
-

7. Conclusion

Three publicly available datasets related to Prompt Injection and Jailbreak Detection — Neuralchemy, Cyberec, and Hlyn — were collected from the Hugging Face Hub and prepared for machine learning through schema standardization, integration, missing-value handling, and duplicate removal. The resulting master dataset of 399,741 prompts was then cleaned and normalized through a full NLP preprocessing pipeline (lowercasing, noise removal, tokenization, stopword removal, and lemmatization), producing a final analysis-ready dataset of 399,741 rows across 11 columns.

Exploratory analysis of the core dataset provided early insight into class balance, attack-category composition, and severity distribution, while the dedicated EDA notebook extends this analysis to the full integrated corpus. The processed dataset is now ready for Feature Engineering, Model Training, Model Evaluation, Hyperparameter Tuning, and Model Deployment as SentinelAI moves into its next development phase.

8. References

- [Neuralchemy Prompt Injection Dataset \(Hugging Face\)](#)
- [Cyberec Prompt Injection Dataset \(Hugging Face\)](#)
- [Hlyn Prompt Injection Judge Dataset \(Hugging Face\)](#)
- [Pandas Documentation](#)
- [NumPy Documentation](#)
- [Matplotlib Documentation](#)
- [Seaborn Documentation](#)
- [NLTK Documentation](#)
- [Scikit-learn Documentation](#)
- [Hugging Face Datasets Documentation](#)